The background of the slide features abstract, overlapping geometric shapes in various shades of blue, ranging from light sky blue to deep navy blue. These shapes are primarily located on the left and right sides of the slide, framing the central text area. The shapes include triangles, polygons, and layered rectangles, creating a modern, architectural feel.

Very Deep Convolutional Networks for Large-Scale Image Recognition

Evandro Dessani

Dez-2017

Intro



- ▶ Convolutional networks (ConvNets) have recently enjoyed a great success in large-scale image and video recognition.
 - ▶ Krizhevsky - 2012;
 - ▶ Zeiler & Fergus - 2013;
 - ▶ Sermanet - 2014; Simonyan & Zisserman - 2014)
- ▶ utilise smaller receptive window size and smaller stride of the first convolutional layer.
- ▶ Another line of improvements dealt with training and testing the networks densely over the whole image and over multiple scales

ConvNet General Premisses

- ▶ Fix other parameters other than depth;
- ▶ Steadily increase the depth of the network by adding more convolutional layers;
- ▶ Use of the very small convolution filters in all layers;



Results

- ▶ significantly more accurate ConvNet architectures;
- ▶ achieve the state-of-the-art accuracy on ILSVRC classification and localisation tasks;
- ▶ applicable to other image recognition datasets;
- ▶ achieve excellent performance even when used as a part of a relatively simple pipelines;





Architecture

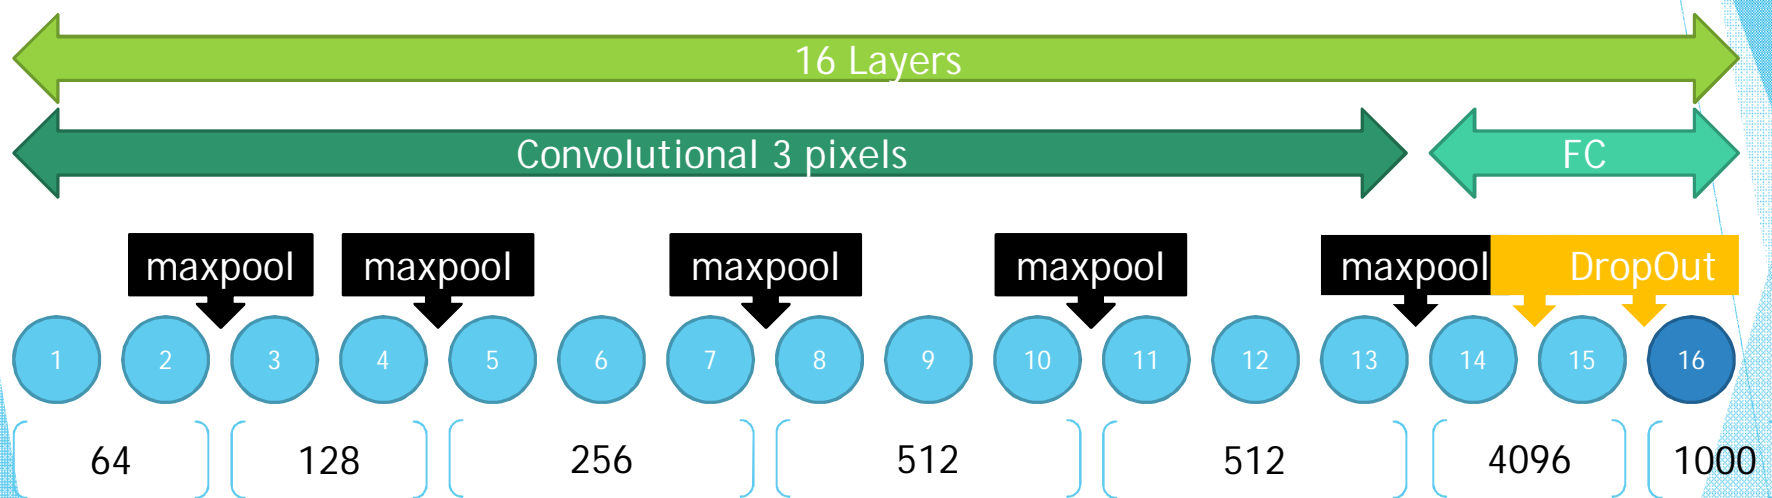
- ▶ The input is as fixed-size 224 x 224 RGB image.
- ▶ 3 x 3 filters on convolutional layers with smallest size for notion of left/right, up/down/center;
- ▶ The convolution stride is fixed to 1 pixel;
- ▶ The padding is 1 pixel for 3×3 conv. Layers;
- ▶ Spatial pooling is carried out by five max-pooling layers followed by some convolutional layers;
- ▶ All hidden layers are equipped with the rectification (ReLU)
- ▶ Max-pooling is performed over a 2×2 pixel window, with stride 2.
- ▶ Three Fully-Connected (FC) layers:
 - ▶ the first two have 4096 channels each, followed by a dropout with 0.5 ratio
 - ▶ the third performs 1000-way ILSVRC classification and thus contains 1000 channels (one for each class).
 - ▶ The final layer is the soft-max layer.

Configurations

- ▶ All configurations follow the generic design and differ only in depth:
 - ▶ From 11 to 19 weight layers where 3FC layers are fixed;
- ▶ The width of conv. layers (the number of channels), start from 64 in the first layer and then increasing by a factor of 2 after each max-pooling layer, until it reaches 512.

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64	conv3-64	conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128	conv3-128	conv3-128
maxpool					
conv3-256	conv3-256	conv3-256	conv3-256	conv3-256	conv3-256
conv3-256	conv3-256	conv3-256	conv1-256	conv3-256	conv3-256
maxpool					
conv3-512	conv3-512	conv3-512	conv3-512	conv3-512	conv3-512
conv3-512	conv3-512	conv3-512	conv1-512	conv3-512	conv3-512
maxpool					
conv3-512	conv3-512	conv3-512	conv3-512	conv3-512	conv3-512
conv3-512	conv3-512	conv3-512	conv1-512	conv3-512	conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

ConvNet Configuration D




● RELU
● SOFTMAX

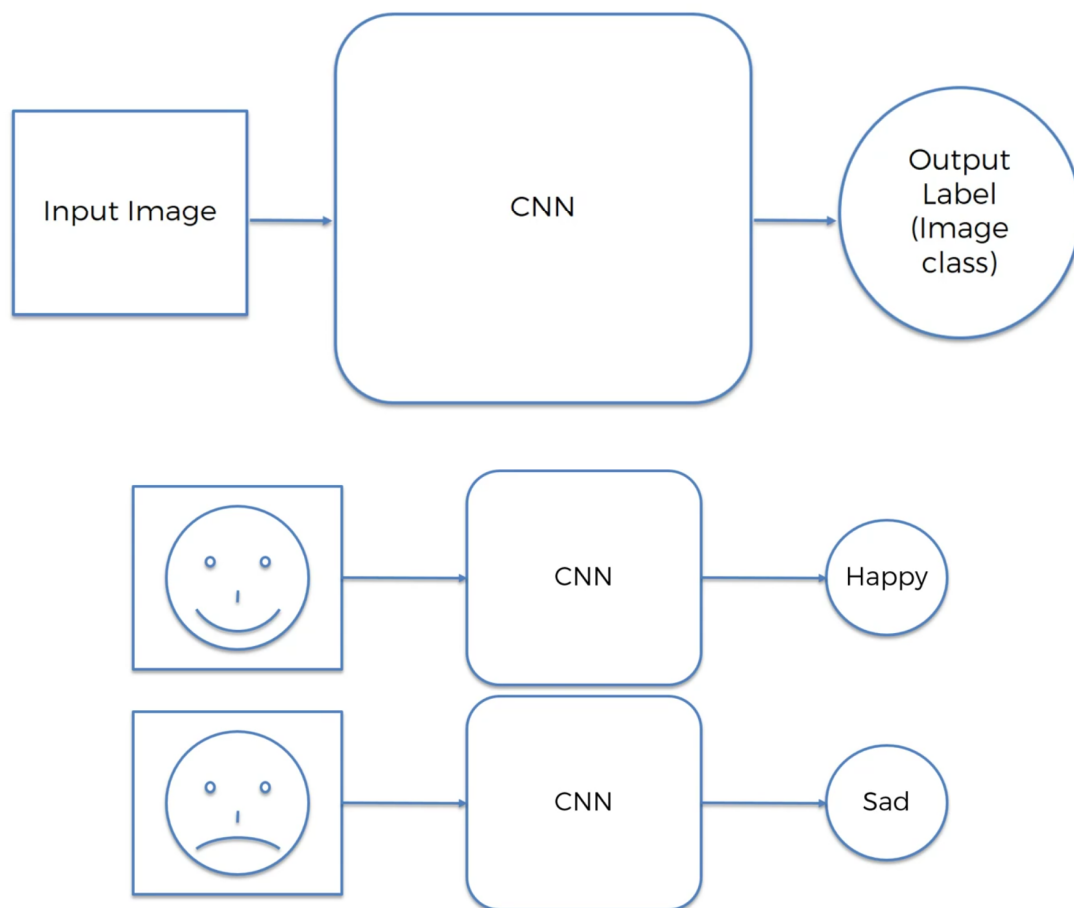
Training




- ▶ Image input sizes = 224 x 224
- ▶ Batch Size = 256
- ▶ Momentum 0.9
- ▶ Dropout Ratio = 0.5
- ▶ Penalty Multiplier = 0.0005
- ▶ Initial Learning Rate = 0.01
- ▶ Final Learning Rate = 0.00001
 - ▶ The learning rate was set to decrease by a factor of 10 when the validation set accuracy stopped improving
 - ▶ In total, the learning rate was decreased 3 times, and the learning was stopped after 370K iterations (74 epochs).
- ▶ Initialization Weights
 - ▶ First four Convolutional Layers and the last three connected layers was initialized with the layers of netconf A
 - ▶ The intermediate layers were initialized randomly

Coding







B / W Image 2x2px



2d array




Colored Image 2x2px





3d array

Red channel

Green channel



Blue channel



0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

STEP 1: Convolution



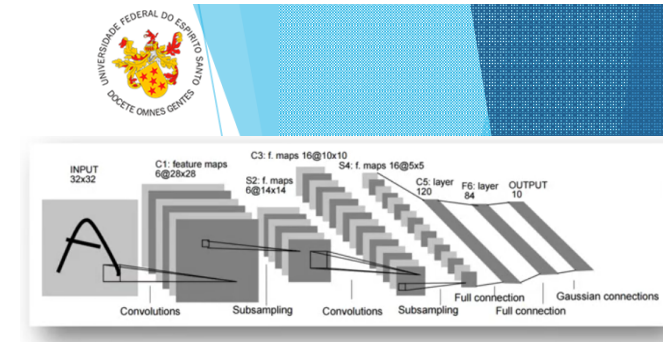
STEP 2: Max Pooling



STEP 3: Flattening



STEP 4: Full Connection



Step 1 - Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

 \otimes

0	0	1
1	0	0
0	1	1

 $=$

0				

Input Image

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Feature Detector

0	0	1
1	0	0
0	1	1

 $=$

0	1	0	0

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

 \otimes

0	0	1
1	0	0
0	1	1

 $=$

0	1	0	0	0

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

 \otimes

0	0	1
1	0	0
0	1	1

 $=$

0	1	0	0	0
0	1			

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

 \otimes

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

 \otimes

0	0	1
1	0	0
0	1	1

 $=$

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4			

Input Image

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

 \otimes

0	0	1
1	0	0
0	1	1

 $=$

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1