

Implantação do README Github LCAD Deep Learning 2017/2

Autores: Renan Mantuanelli de Aquino e Deivison Augusto da Vitória

Passo a passo de como criar uma CNN (Convolutional Neural Network) para identificação de carro, gato e cachorro usando Keras e Tensor Flow baseado na AlexNet e ZFNet usando Windows 10.

1. Instalar anaconda phyton
 - <https://www.continuum.io/downloads>)
2. Instalar spyder
 - Abrir anaconda navegador e instalar spyder
3. Instalar theano
 - Abrir anaconda prompt e instalar theano: pip install theano
4. Instalar tensorflow
 - Abrir anaconda prompt e instalar theano: pip install tensorflow
5. Instalar keras
 - Abrir anaconda prompt e instalar keras: pip install keras
6. Atualizar todos os programas recém instalados e outros pacotes do Anaconda
 - Abrir anaconda prompt e atualizar: conda update -all
7. Instalar o Vstudio 2015
8. Instalar cuda_8.0.61_win10
9. Instalar cuda_8.0.61.2_windows
10. Incluir caminho no Path:
 - edit system variables
 - Environment variables
 - System variables
 - Path Edit
 - New (inserir na última linha caminho que está instalado o cuda ex.: C:\Program Files\NVIDIA GPU Computing Toolkit\CUDA\bin)
11. Instalar GPU para utilização no Tensorflow
 - https://www.tensorflow.org/install/install_windows
 - Criar ambiente conda com nome tensorflow (abrir anaconda prompt)
 - ✓ conda create -n tensorflow python=3.5
 - Ativar ambiente conda (abrir anaconda prompt)
 - ✓ activate tensorflow
 - Atualizar tensorflow (abrir anaconda prompt)
 - ✓ pip install --ignore-installed --upgrade tensorflow
 - Instalar GPU versão do tensorflow
 - ✓ pip install --ignore-installed --upgrade tensorflow-gpu
 - Testar sucesso instalação digitando os comandos abaixo no spyder
 - ✓ Abrir spyder (phynton)
 - ✓ import tensorflow as tf
 - hello = tf.constant('Hello, TensorFlow!')
 - sess = tf.Session()
 - print(sess.run(hello))
 - Se a instalação foi concluída com sucesso a saída (Console Phyton - Spyder) irá retornar *Hello, TensorFlow!*
12. Criar a rede CNN baseado na ZFnet


"""

Spyder Editor

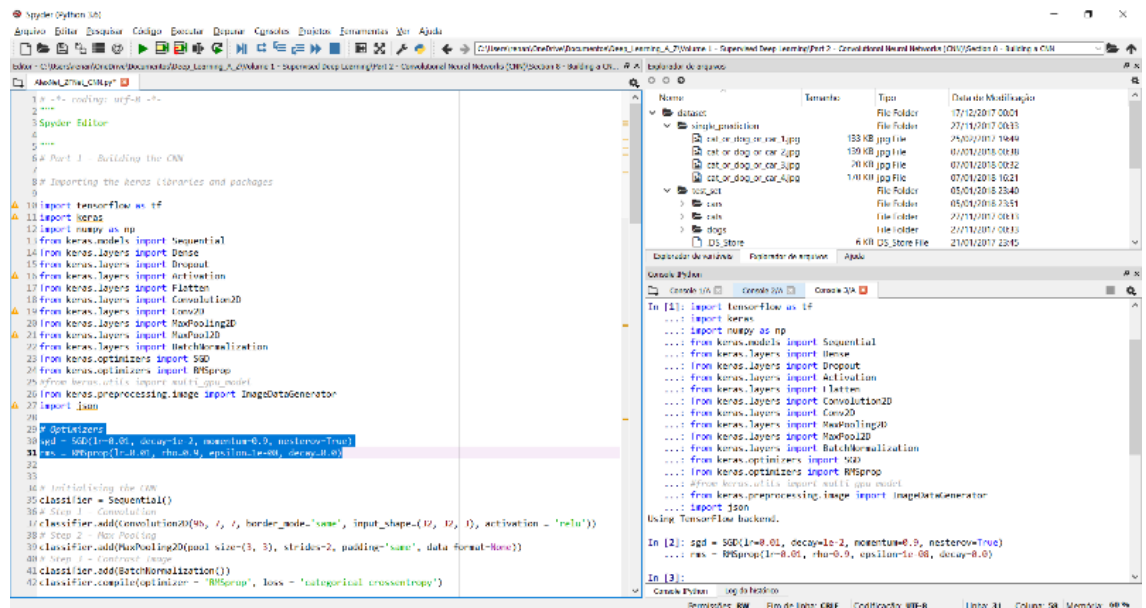
"""

Part 1 – Construindo a rede Aquino&VitóriaNet (CNN) baseado na ZFNET

```
import tensorflow as tf
import keras
import numpy as np
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Dropout
from keras.layers import Activation
from keras.layers import Flatten
from keras.layers import Convolution2D
from keras.layers import Conv2D
from keras.layers import MaxPooling2D
from keras.layers import MaxPool2D
from keras.layers import BatchNormalization
from keras.optimizers import SGD
from keras.optimizers import RMSprop
from keras.utils import multi_gpu_model
from keras.preprocessing.image import ImageDataGenerator
import json
```



```
rms = RMSprop(lr=0.01, rho=0.9, epsilon=1e-08, decay=0.0)
```



Inicializando a rede Aquino&VitóriaFNET CNN

O tipo da rede será sequencial que é uma característica da CNN

classifier = Sequential()

Step 1 – Convolution – 96 camadas, kernel 7x7, imagem entrada 32x32x3(RGB), função de ativação ReLU

classifier.add(Convolution2D(96, 7, 7, border_mode='same', input_shape=(32, 32, 3), activation = 'relu'))

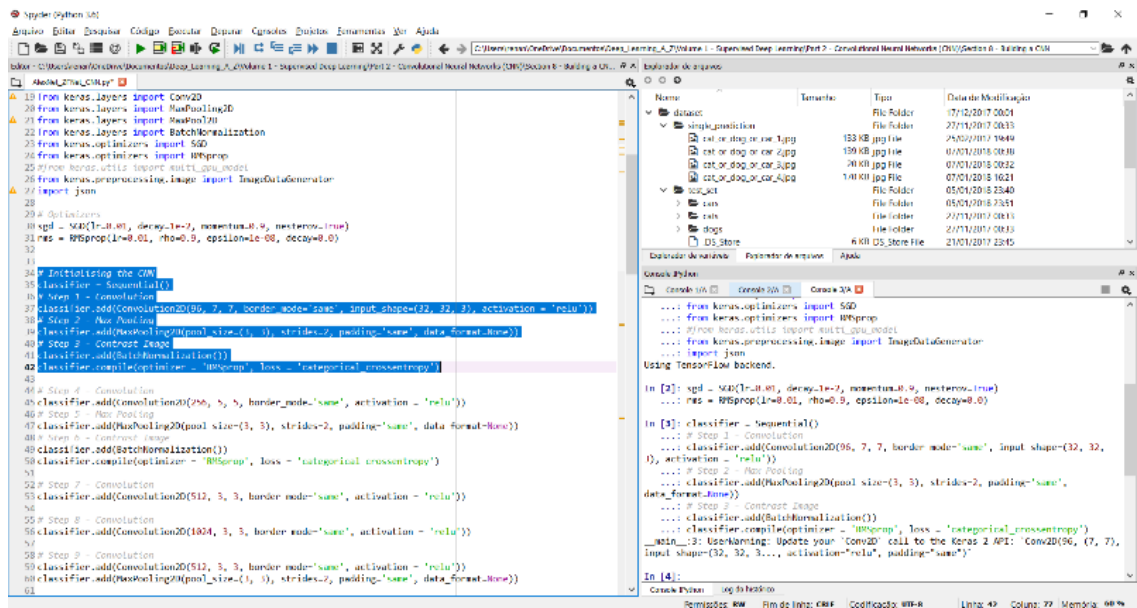
Step 2 - Max Pooling – tamanho 3x3, passo2

classifier.add(MaxPooling2D(pool_size=(3, 3), strides=2, padding='same', data_format=None))

Step 3 - Contrast Image – aplicando operação de contraste local de normalização na saída do mapa de caracterísica (feature map)

classifier.add(BatchNormalization())

classifier.compile(optimizer = 'RMSprop', loss = 'categorical_crossentropy')



Step 4 – Convolution - 256 camadas, kernel 5x5, imagem entrada 16x16x256 função de ativação ReLU

classifier.add(Convolution2D(256, 5, 5, border_mode='same', activation = 'relu'))

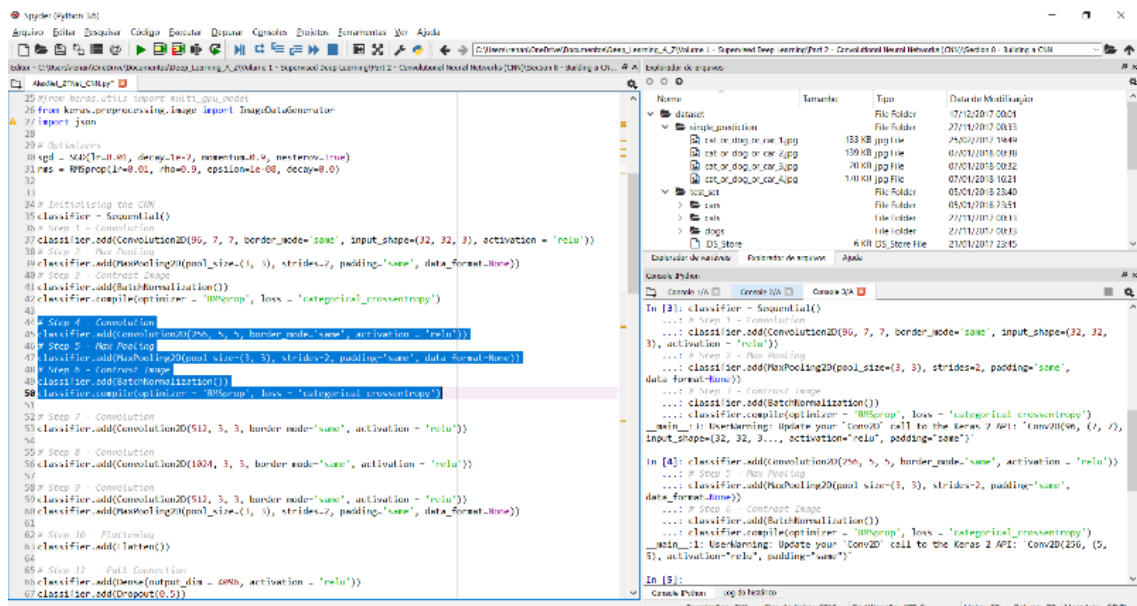
Step 5 - Max Pooling - tamanho 3x3, passo2

classifier.add(MaxPooling2D(pool_size=(3, 3), strides=2, padding='same', data_format=None))

Step 6 - Contrast Image - aplicando operação de contraste local de normalização na saída do mapa de característica (feature map)

classifier.add(BatchNormalization())

classifier.compile(optimizer = 'RMSprop', loss = 'categorical_crossentropy')



Step 7 – Convolution - 512 camadas, kernel 3x3, imagem entrada 8x8x512, função de ativação ReLU

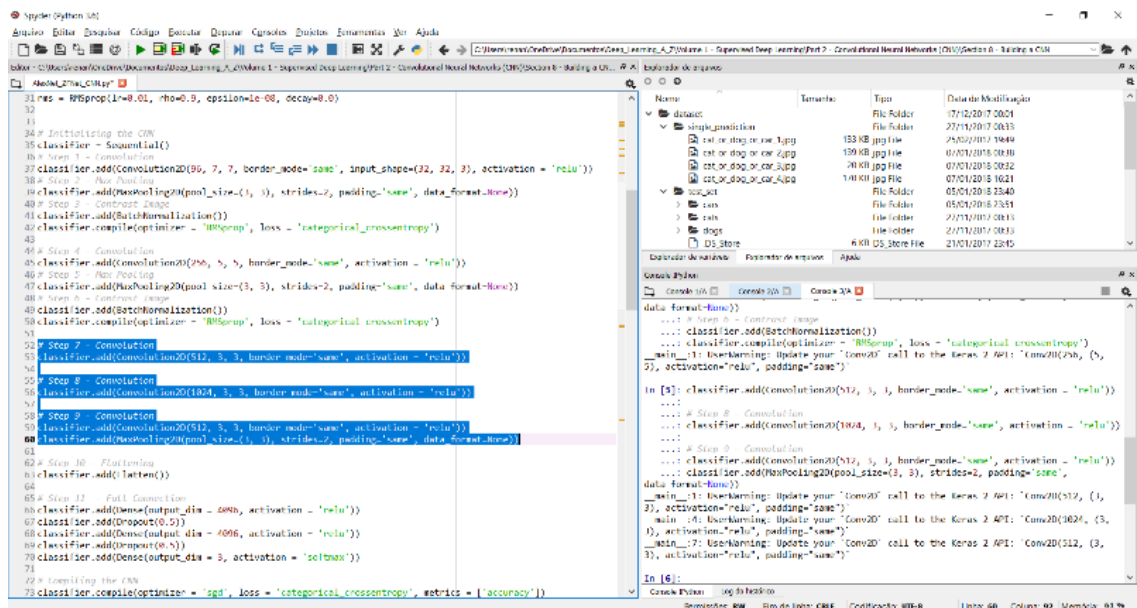
```
classifier.add(Convolution2D(512, 3, 3, border_mode='same', activation = 'relu'))
```

Step 8 - Convolution - 1024 camadas, kernel 3x3, imagem entrada 8x8x1024, função de ativação ReLU

```
classifier.add(Convolution2D(1024, 3, 3, border_mode='same', activation = 'relu'))
```

Step 9 - Convolution - 512 camadas, kernel 3x3, imagem entrada 8x8x512, função de ativação ReLU

```
classifier.add(Convolution2D(512, 3, 3, border_mode='same', activation = 'relu'))
```



```
31 m2 = RMSprop(lr=0.01, rho=0.9, epsilon=1e-08, decay=0.0)
32
33 # Step 7 - Convolution
34 # Step 8 - Convolution
35 classifier.add(Convolution2D(1024, 3, 3, border_mode='same', input_shape=(32, 32, 3), activation = 'relu'))
36 # Step 9 - Convolution
37 classifier.add(Convolution2D(512, 3, 3, border_mode='same', input_shape=(32, 32, 3), activation = 'relu'))
38 # Step 10 - Max Pooling
39 classifier.add(MaxPooling2D(pool_size=(3, 3), strides=2, padding='same', data_format='none'))
40 # Step 11 - Flatten
41 classifier.add(Flatten())
42 classifier.compile(optimizer = 'RMSprop', loss = 'categorical_crossentropy')
43
44 # Step 12 - Full Connection
45 classifier.add(Dense(output_dim=4096, activation = 'relu'))
46 classifier.add(Dense(output_dim=4096, activation = 'relu'))
47 classifier.add(Dense(output_dim=1000, activation = 'softmax'))
48 classifier.compile(optimizer = 'RMSprop', loss = 'categorical_crossentropy', metrics = ['accuracy'])
49
50 # Step 13 - Training
51 classifier.fit_generator(generator, nb_epoch=10, validation_data=(validation_data, validation_labels))
52
53 # Step 14 - Evaluation
54 classifier.evaluate_generator(generator, nb_epoch=10)
```

Step 10 - Max Pooling - tamanho 3x3, passo2

```
classifier.add(MaxPooling2D(pool_size=(3, 3), strides=2, padding='same', data_format='none'))
```

Step 11 – Flattening – preparação da saída da última etapa de convolução conforme necessário para entrada na camada primeira totalmente conectada com 8192 saídas (dimensões).

```
classifier.add(Flatten())
```

Step 12 - Full Connection – primeira camada totalmente conectada com 4096 dimensões, 33.558.528 parâmetros e função de ativação ReLU, desligamento de 50% dos neurônios (dropout). A segunda camada totalmente conectada com 4096 dimensões, 16.781.312 parâmetros e função de ativação ReLU, desligamento de 50% dos neurônios (dropout). A última camada classificadora do tipo softmax com três saídas (carro, gato ou cachorro) com 12.291 parâmetros.

No caso da ZFNet a saída da última camada seria 1000, em funções da possibilidade de classificação de 1000 classes diferentes.

```

classifier.add(Dense(output_dim = 4096, activation = 'relu'))
classifier.add(Dropout(0.5))
classifier.add(Dense(output_dim = 4096, activation = 'relu'))
classifier.add(Dropout(0.5))
classifier.add(Dense(output_dim = 3, activation = 'softmax'))

```

```

46 # Step 5 - Pooling
47 classifier.add(MaxPooling2D(pool_size=(3, 3), strides=2, padding='same', data_format='none'))
48 # Step 6 - Convolution
49 classifier.add(BatchNormalization())
50 classifier.compile(optimizer = 'RMSprop', loss = 'categorical_crossentropy')
51
52 # Step 7 - Convolution
53 classifier.add(Convolution2D(512, 3, 3, border_mode='same', activation = 'relu'))
54
55 # Step 8 - Convolution
56 classifier.add(Convolution2D(1024, 3, 3, border_mode='same', activation = 'relu'))
57
58 # Step 9 - Convolution
59 classifier.add(Convolution2D(512, 3, 3, border_mode='same', activation = 'relu'))
60 classifier.add(MaxPooling2D(pool_size=(1, 3), strides=2, padding='same', data_format='none'))
61
62 # Step 10 - Flatten
63 classifier.add(Flatten())
64
65 # Step 12 - Full Connection
66 classifier.add(Dense(output_dim = 4096, activation = 'relu'))
67 classifier.add(Dropout(0.5))
68 classifier.add(Dense(output_dim = 4096, activation = 'relu'))
69 classifier.add(Dropout(0.5))
70 classifier.add(Dense(output_dim = 3, activation = 'softmax'))
71
72 # Compiling the CNN
73 classifier.compile(optimizer = 'sgd', loss = 'categorical_crossentropy', metrics = ['accuracy'])
74
75 classifier.summary()
76
77 # Part 2 - Fitting the CNN to the images
78 from keras.preprocessing.image import ImageDataGenerator
79
80 train_datagen = ImageDataGenerator(
81     rotation=1.25,
82     zoom_range=0.2,
83     shear_range=0.2,
84     horizontal_flip=True)
85
86 test_datagen = ImageDataGenerator(rotation=1.25)
87
88 training_set = train_datagen.flow_from_directory('dataset/training_set',
89
90

```

Compiling the CNN – Compilando a rede neural para classificação final (acuracidade) e verificação da arquitetura da rede

```

classifier.compile(optimizer = 'sgd', loss = 'categorical_crossentropy', metrics = ['accuracy'])

```

```

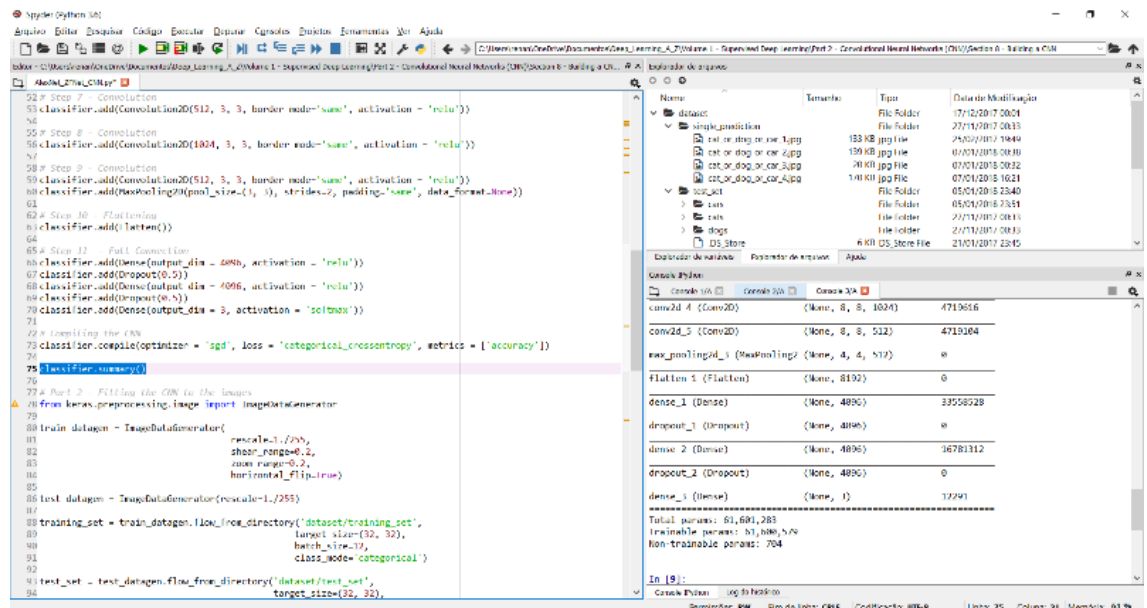
52 # Step 7 - Convolution
53 classifier.add(Convolution2D(512, 3, 3, border_mode='same', activation = 'relu'))
54
55 # Step 8 - Convolution
56 classifier.add(Convolution2D(1024, 3, 3, border_mode='same', activation = 'relu'))
57
58 # Step 9 - Convolution
59 classifier.add(Convolution2D(512, 3, 3, border_mode='same', activation = 'relu'))
60 classifier.add(MaxPooling2D(pool_size=(1, 3), strides=2, padding='same', data_format='none'))
61
62 # Step 10 - Flatten
63 classifier.add(Flatten())
64
65 # Step 12 - Full Connection
66 classifier.add(Dense(output_dim = 4096, activation = 'relu'))
67 classifier.add(Dropout(0.5))
68 classifier.add(Dense(output_dim = 4096, activation = 'relu'))
69 classifier.add(Dropout(0.5))
70 classifier.add(Dense(output_dim = 3, activation = 'softmax'))
71
72 # Compiling the CNN
73 classifier.compile(optimizer = 'sgd', loss = 'categorical_crossentropy', metrics = ['accuracy'])
74
75 classifier.summary()
76
77 # Part 2 - Fitting the CNN to the images
78 from keras.preprocessing.image import ImageDataGenerator
79
80 train_datagen = ImageDataGenerator(
81     rotation=1.25,
82     zoom_range=0.2,
83     shear_range=0.2,
84     horizontal_flip=True)
85
86 test_datagen = ImageDataGenerator(rotation=1.25)
87
88 training_set = train_datagen.flow_from_directory('dataset/training_set',
89     target_size=(32, 32),
90     batch_size=32,
91     class_mode='categorical')
92
93 test_set = test_datagen.flow_from_directory('dataset/test_set',
94     target_size=(32, 32),

```

```

classifier.summary()

```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 96)	14208
max_pooling2d_1 (MaxPooling2D)	(None, 16, 16, 96)	0
batch_normalization_1 (Batch Normalization)	(None, 16, 16, 96)	384
conv2d_2 (Conv2D)	(None, 16, 16, 256)	614656
max_pooling2d_2 (MaxPooling2D)	(None, 8, 8, 256)	0
batch_normalization_2 (Batch Normalization)	(None, 8, 8, 256)	1024
conv2d_3 (Conv2D)	(None, 8, 8, 512)	1180160
conv2d_4 (Conv2D)	(None, 8, 8, 1024)	4719616
conv2d_5 (Conv2D)	(None, 8, 8, 512)	4719104
max_pooling2d_3 (MaxPooling2D)	(None, 4, 4, 512)	0
flatten_1 (Flatten)	(None, 8192)	0
dense_1 (Dense)	(None, 4096)	33558528
dropout_1 (Dropout)	(None, 4096)	0
dense_2 (Dense)	(None, 4096)	16781312
dropout_2 (Dropout)	(None, 4096)	0
dense_3 (Dense)	(None, 3)	12291
Total params: 61,601,283		
Trainable params: 61,600,579		
Non-trainable params: 704		

Part 2 – Fazendo a ligação da rede neural com as imagens para treino, teste/validação

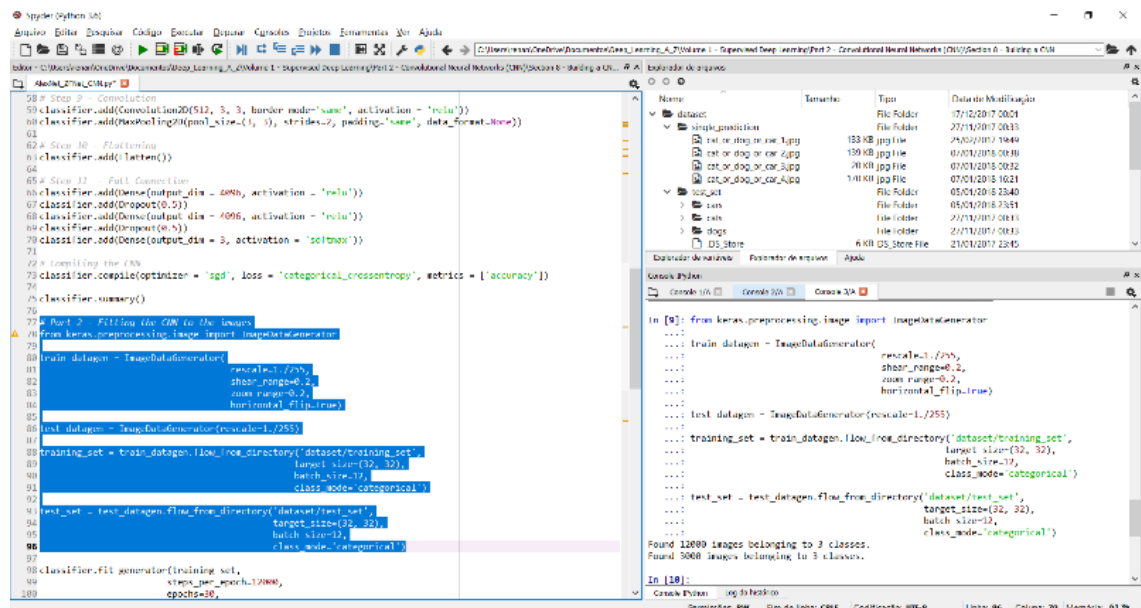
```
from keras.preprocessing.image import ImageDataGenerator
```

```
train_datagen = ImageDataGenerator(
    rescale=1./255,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True)
```

```
test_datagen = ImageDataGenerator(rescale=1./255)
```

```
training_set = train_datagen.flow_from_directory('dataset/training_set',
    target_size=(32, 32),
    batch_size=12,
    class_mode='categorical')
```

```
test_set = test_datagen.flow_from_directory('dataset/test_set',
    target_size=(32, 32),
    batch_size=12,
    class_mode='categorical')
```



```
classifier.fit_generator(training_set,
    steps_per_epoch=12000,
    epochs=30,
    validation_data=test_set,
    validation_steps=3000)
```

Part 3 – Salvando os pesos após treino (aprendizado “memória” da rede)

```
classifier.save("TesteFinal.h5")
print('Classifier Saved')
```



```

68 classifier.add(Dense(output_dim = 4096, activation = 'relu'))
69 classifier.add(Dense(output_dim = 5, activation = 'softmax'))
70 classifier.compile(optimizer = 'sgd', loss = 'categorical_crossentropy', metrics = ['accuracy'])
71
72 # Creating the CNN
73 classifier.compile(optimizer = 'sgd', loss = 'categorical_crossentropy', metrics = ['accuracy'])
74
75 classifier.summary()
76
77 # Part 2 - Fitting the CNN to the images
78 from keras.preprocessing.image import ImageDataGenerator
79
80 train_datagen = ImageDataGenerator(
81     rescale=1./255,
82     shear_range=0.2,
83     zoom_range=0.2,
84     horizontal_flip=True)
85
86 test_datagen = ImageDataGenerator(rescale=1./255)
87
88 training_set = train_datagen.flow_from_directory('dataset/training_set',
89     target_size=(32, 32),
90     batch_size=32,
91     class_mode='categorical')
92
93 test_set = test_datagen.flow_from_directory('dataset/test_set',
94     target_size=(32, 32),
95     batch_size=32,
96     class_mode='categorical')
97
98 classifier.fit_generator(training_set,
99     steps_per_epoch=12000,
100     epochs=10,
101     validation_data=test_set,
102     validation_steps=3000)
103
104 # Part 3 - Save the weights
105 classifier.save('TestFinal.h5')
106 print('Classifier Saved')
107
108 # Part 4 - Making new predictions
109 import numpy as np
110 from keras.preprocessing import image

```

dataset

- single_prediction
 - cat_or_dog_or_car_1.jpg
 - cat_or_dog_or_car_2.jpg
 - cat_or_dog_or_car_3.jpg
 - cat_or_dog_or_car_4.jpg
- test_set
 - cars
 - cats
 - dogs
 - DS_Store

dataset/training_set

- cars
- cats
- dogs
- DS_Store

dataset/test_set

- cars
- cats
- dogs
- DS_Store

In [10]:

```

classifier.fit_generator(training_set,
    steps_per_epoch=12000,
    epochs=10,
    validation_data=test_set,
    validation_steps=3000)

```

In [10]:

```

classifier.save('TestFinal.h5')
print('Classifier Saved')

```

In [10]:

```

import numpy as np
from keras.preprocessing import image

```

Part 4 – Fazendo novas predições com a rede já treinada

import numpy as np
from keras.preprocessing import image

```

82 shear_range=0.2,
83 zoom_range=0.2,
84 horizontal_flip=True)
85
86 test_datagen = ImageDataGenerator(rescale=1./255)
87
88 training_set = train_datagen.flow_from_directory('dataset/training_set',
89     target_size=(32, 32),
90     batch_size=32,
91     class_mode='categorical')
92
93 test_set = test_datagen.flow_from_directory('dataset/test_set',
94     target_size=(32, 32),
95     batch_size=32,
96     class_mode='categorical')
97
98 classifier.fit_generator(training_set,
99     steps_per_epoch=12000,
100     epochs=10,
101     validation_data=test_set,
102     validation_steps=3000)
103
104 # Part 3 - Save the weights
105 classifier.save('TestFinal.h5')
106 print('Classifier Saved')
107
108 # Part 4 - Making new predictions
109 import numpy as np
110 from keras.preprocessing import image
111
112 test_image = image.load_img('dataset/single_prediction/cat_or_dog_or_car_3.jpg', target_size = (32, 32))
113 test_image = image.img_to_array(test_image)
114 test_image = np.expand_dims(test_image, axis = 0)
115 result = classifier.predict(test_image)
116 print(result)
117 training_set.class_indices
118
119 test_image = image.load_img('dataset/single_prediction/cat_or_dog_or_car_2.jpg', target_size = (32, 32))
120 test_image = image.img_to_array(test_image)
121 test_image = np.expand_dims(test_image, axis = 0)
122 result = classifier.predict(test_image)
123 print(result)
124 training_set.class_indices

```

dataset

- single_prediction
 - cat_or_dog_or_car_1.jpg
 - cat_or_dog_or_car_2.jpg
 - cat_or_dog_or_car_3.jpg
 - cat_or_dog_or_car_4.jpg
- test_set
 - cars
 - cats
 - dogs
 - DS_Store

dataset/training_set

- cars
- cats
- dogs
- DS_Store

dataset/test_set

- cars
- cats
- dogs
- DS_Store

In [14]:

```

import numpy as np
from keras.preprocessing import image

```

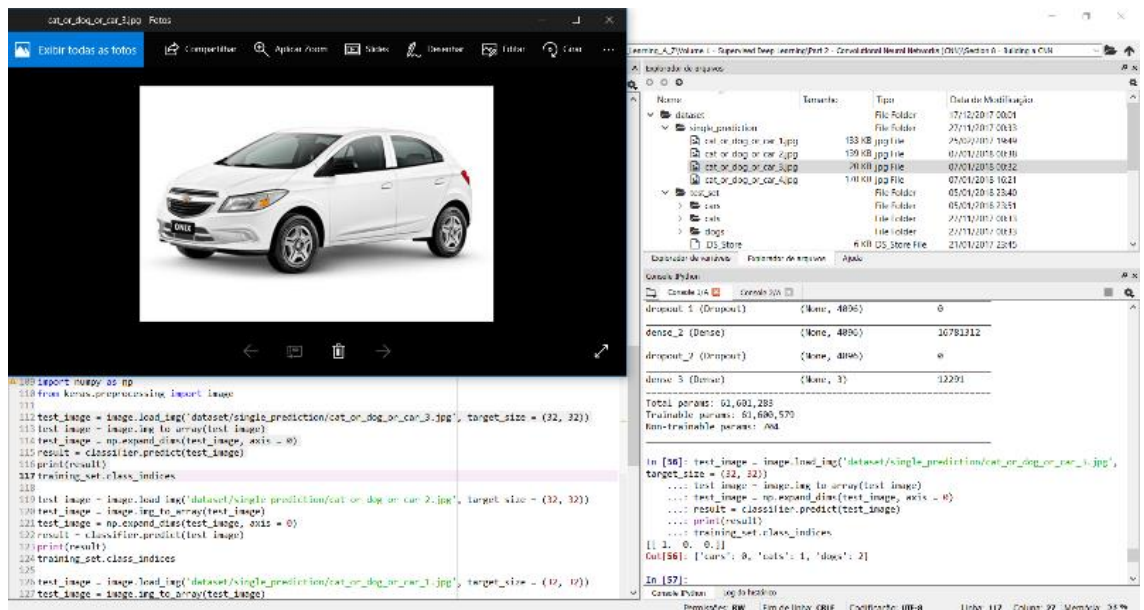
In [15]:

```

test_image = image.load_img('dataset/single_prediction/cat_or_dog_or_car_3.jpg', target_size = (32, 32))
test_image = image.img_to_array(test_image)
test_image = np.expand_dims(test_image, axis = 0)
result = classifier.predict(test_image)
print(result)
training_set.class_indices

```

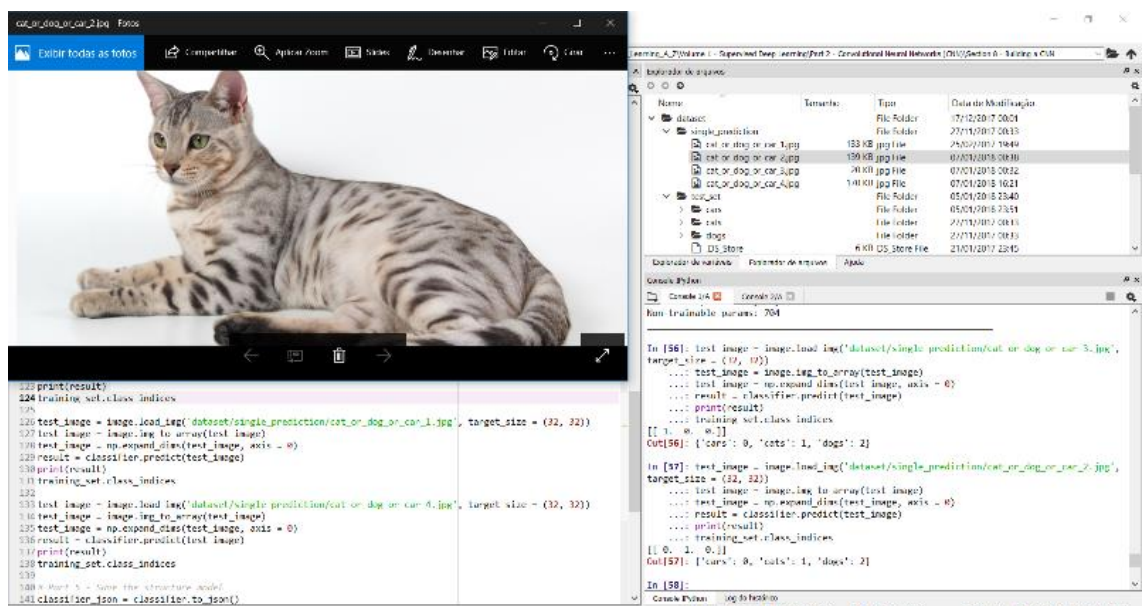
test_image = image.load_img('dataset/single_prediction/cat_or_dog_or_car_3.jpg',
target_size = (32, 32))
test_image = image.img_to_array(test_image)
test_image = np.expand_dims(test_image, axis = 0)
result = classifier.predict(test_image)
print(result)
training_set.class_indices



```

test_image = image.load_img('dataset/single_prediction/cat_or_dog_or_car_2.jpg',
target_size = (32, 32))
test_image = image.img_to_array(test_image)
test_image = np.expand_dims(test_image, axis = 0)
result = classifier.predict(test_image)
print(result)
training_set.class_indices

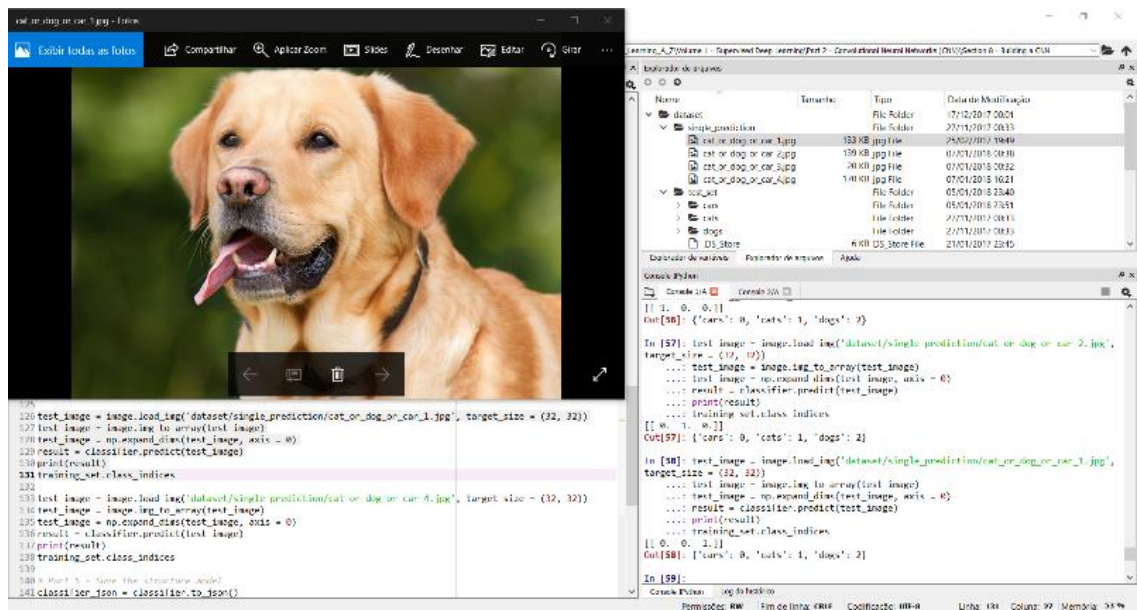
```



```

test_image = image.load_img('dataset/single_prediction/cat_or_dog_or_car_1.jpg',
target_size = (32, 32))
test_image = image.img_to_array(test_image)
test_image = np.expand_dims(test_image, axis = 0)
result = classifier.predict(test_image)
print(result)
training_set.class_indices

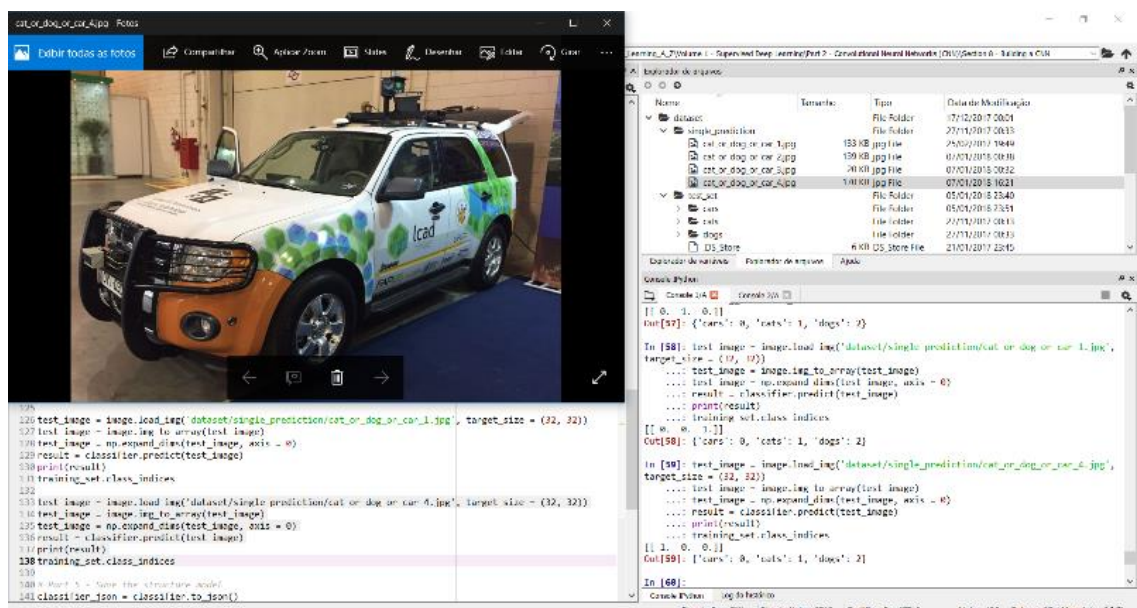
```



```

test_image = image.load_img('dataset/single_prediction/cat_or_dog_or_car_4.jpg',
target_size = (32, 32))
test_image = image.img_to_array(test_image)
test_image = np.expand_dims(test_image, axis = 0)
result = classifier.predict(test_image)
print(result)
training_set.class_indices

```



Part 5 – Salvando a estrutura do modelo

```

classifier_json = classifier.to_json()
with open("Aquino&VitóriaNetStructure.json", "w") as json_file:
    json_file.write(classifier_json)

```

